

FRI and Proximity Proofs: what they are, what they are for, and the future

(Part 2)

Dan Boneh
Stanford University

FRI: a Reed-Solomon IOP of Proximity

Let $\mathcal{C} = \text{RS}[\mathbb{F}, \mathcal{L}, d]$, $u_0: \mathcal{L} \rightarrow \mathbb{F}$, and $\delta \in [0,1]$

Prover $P(\mathcal{C}, u_0, \cdot)$

Verifier $V^{u_0}(\mathcal{C})$

Goal:

- $u_0 \in \text{RS}[\mathbb{F}, \mathcal{L}, d] \Rightarrow$ Verifier outputs accept (with prob. 1)
- u_0 is δ -far from $\text{RS}[\mathbb{F}, \mathcal{L}, d] \Rightarrow$ Verifier outputs reject w.h.p

We don't care what happens when u_0 is between the two cases

Why needed? A key tool for compiling a Poly-IOP into an IOP.

FRI: a Reed-Solomon IOP of Proximity

Let $\mathcal{C} = \text{RS}[\mathbb{F}, \mathcal{L}, d]$, $u_0: \mathcal{L} \rightarrow \mathbb{F}$, and $\delta \in [0,1]$

Prover $P(\mathcal{C}, u_0, \cdot)$

Verifier $V^{u_0}(\mathcal{C})$

Recall that:

$\mathcal{L} = \{1, \omega, \omega^2, \dots, \omega^{n-1}\} \subseteq \mathbb{F}$ and n, d are powers of two,
where $\omega^n = 1$ is an n -th primitive root of unity

Then: $|\mathcal{L}^2| := |\{a^2: a \in \mathcal{L}\}| = |\mathcal{L}|/2 = n/2$ ($-a, a \rightarrow a^2$)

$|\mathcal{L}^4| := |\{a^4: a \in \mathcal{L}\}| = |\mathcal{L}|/4 = n/4$

Review: 2-way folding an arbitrary word $u: \mathcal{L} \rightarrow \mathbb{F}$

For $u: \mathcal{L} \rightarrow \mathbb{F}$ and $r \in \mathbb{F}$ define $u_e, u_o, u_{\text{fold},r}: \mathcal{L}^2 \rightarrow \mathbb{F}$ as

- for $a \in \mathcal{L}$: $u_e(a^2) := \frac{u(a)+u(-a)}{2}$ and $u_o(a^2) := \frac{u(a)-u(-a)}{2a}$
- for $b \in \mathcal{L}^2$: $u_{\text{fold},r}(b) := u_e(b) + r \cdot u_o(b)$

Can similarly define 2^w -way folding for $w \geq 1$.

Recall $|\mathcal{L}^2| = |\mathcal{L}|/2$

Review: an important corollary

Let $\mathcal{C} = \text{RS}[\mathbb{F}, \mathcal{L}, d]$ and $\mathcal{C}' = \text{RS}[\mathbb{F}, \mathcal{L}^2, d/2]$

Corollary: For every $u: \mathcal{L} \rightarrow \mathbb{F}$

(folding does not decrease distance, w.h.p)

- if $\Delta(u, \mathcal{C}) < 1 - \sqrt{\rho}$ then $\Pr_r[\Delta(u_{\text{fold}, r}, \mathcal{C}') \geq \Delta(u, \mathcal{C})] \geq 1 - \text{err}$
- if $\Delta(u, \mathcal{C}) \geq 1 - \sqrt{\rho}$ then $\Pr_r[\Delta(u_{\text{fold}, r}, \mathcal{C}') \geq 1 - \sqrt{\rho}] \geq 1 - \text{err}$

How FRI works

A Reed-Solomon IOP of Proximity (RS-IOPP)

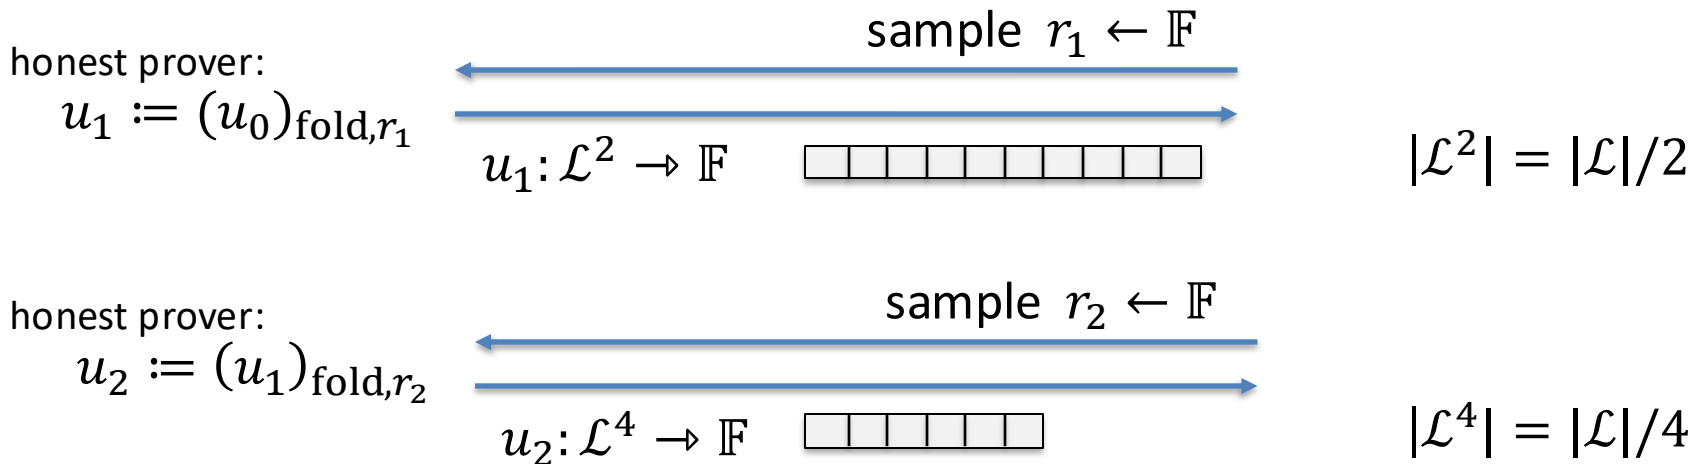
FRI phase 1: commit phase

Let $\mathcal{C} = \text{RS}[\mathbb{F}, \mathcal{L}, d]$, $u_0: \mathcal{L} \rightarrow \mathbb{F}$, and $\delta \in [0, 1]$

Prover $P(\mathcal{C}, u_0, \cdot)$

Verifier $V^{u_0}(\mathcal{C})$

Phase 1: (commit)





FRI phase 1: commit phase

Let $\mathcal{C} = \text{RS}[\mathbb{F}, \mathcal{L}, d]$, $u_0: \mathcal{L} \rightarrow \mathbb{F}$, and $\delta \in [0, 1]$

Prover $P(\mathcal{C}, u_0, \cdot)$

Verifier $V^{u_0}(\mathcal{C})$

Phase 1: (commit)



$$|\mathcal{L}^{2^t}| = |\mathcal{L}|/2^t$$

FRI phase 2: query phase

Let $\mathcal{C} = \text{RS}[\mathbb{F}, \mathcal{L}, d]$, $u_0: \mathcal{L} \rightarrow \mathbb{F}$, and $\delta \in [0, 1]$


Phase 2: (query)

Verifier $V^{u_0, \dots, u_t}(\mathcal{C}, r_1, \dots, r_t, u_t)$

$u_1: \mathcal{L}^2 \rightarrow \mathbb{F}$ 

$u_2: \mathcal{L}^4 \rightarrow \mathbb{F}$ 

\vdots \vdots

$u_t: \mathcal{L}^{2^t} \rightarrow \mathbb{F}$ 

for $i = 1, 2, \dots, t$:

spot check that $u_i = (u_{i-1})_{\text{fold}, r_i}$

output yes if $u_t \in \text{RS}[\mathbb{F}, \mathcal{L}^{2^t}, d/2^t]$

[Prover sent Merkle commits to u_1, \dots, u_t]

Note: prover sends short u_t to verifier explicitly

(FRI terminates when u_t is “short enough”)

FRI phase 2: query phase

Let $\mathcal{C} = \text{RS}[\mathbb{F}, \mathcal{L}, d]$, $u_0: \mathcal{L} \rightarrow \mathbb{F}$, and $\delta \in [0,1]$


Phase 2: (query)

Verifier $V^{u_0, \dots, u_t}(\mathcal{C}, r_1, \dots, r_t, u_t)$

$u_1: \mathcal{L}^2 \rightarrow \mathbb{F}$ 

$u_2: \mathcal{L}^4 \rightarrow \mathbb{F}$ 

\vdots

$u_t: \mathcal{L}^{2^t} \rightarrow \mathbb{F}$ 

for $i = 1, 2, \dots, t$:

spot check that $u_i = (u_{i-1})_{\text{fold}, r_i}$

output yes if $u_t \in \text{RS}[\mathbb{F}, \mathcal{L}^{2^t}, d/2^t]$

Why is this δ -sound? Intuition: u_0 is δ -far from $\text{RS}[\mathbb{F}, \mathcal{L}, d] \Rightarrow$

u_1 is “far” from $\text{RS}[\mathbb{F}, \mathcal{L}^2, d/2] \Rightarrow \dots \Rightarrow u_t$ is “far” from $\text{RS}[\mathbb{F}, \mathcal{L}^{2^t}, d/2^t]$

How to spot check: method 1

Let $\mathcal{C} = \text{RS}[\mathbb{F}, \mathcal{L}, d]$, $u_0: \mathcal{L} \rightarrow \mathbb{F}$, and $\delta \in [0,1]$


Phase 2: (query)

Verifier $V^{u_0, \dots, u_t}(\mathcal{C}, r_1, \dots, r_t, u_t)$

$u_1: \mathcal{L}^2 \rightarrow \mathbb{F}$ 

$u_2: \mathcal{L}^4 \rightarrow \mathbb{F}$ 

\vdots

$u_t: \mathcal{L}^{2^t} \rightarrow \mathbb{F}$ 

$$z = \frac{u_{i-1}(s) + u_{i-1}(-s)}{2} + r_i \cdot \frac{u_{i-1}(s) - u_{i-1}(-s)}{2s}$$

How to check that $u_i = (u_{i-1})_{\text{fold}, r_i}$:

Repeat m times:

- choose random $s \in \mathcal{L}^{2^{i-1}}$
- query $u_{i-1}(s)$, $u_{i-1}(-s)$, $u_i(s^2)$
- compute $z := (u_{i-1})_{\text{fold}, r_i}(s^2)$
- reject if $z \neq u_i(s^2)$

How to spot check in a picture


$m = 2$: (spot check at two random spots per oracle)

$u_0: \mathcal{L} \rightarrow \mathbb{F}$

$u_1: \mathcal{L}^2 \rightarrow \mathbb{F}$

$u_2: \mathcal{L}^4 \rightarrow \mathbb{F}$

$u_3: \mathcal{L}^8 \rightarrow \mathbb{F}$



$$u_1(s^2) \stackrel{?}{=} \frac{u_0(s) + u_0(-s)}{2} + r_i \cdot \frac{u_0(s) - u_0(-s)}{2s}$$

Since prover opens $(s, -s)$ jointly,
it places both into a single leaf
of Merkle tree

\Rightarrow one query to get both

Reject if any spot checks fail

Total of $2mt$ queries to oracles: $2m$ per inner oracle (u_1, u_2).

Why is this protocol sound?

For $i = 0, \dots, t$: let

- $\mathcal{C}_i := \text{RS}[\mathbb{F}, \mathcal{L}^{2^i}, d/2^i]$ and
- $\eta_i := (\text{distance of } u_i \text{ to } \mathcal{C}_i) = \Delta(u_i, \mathcal{C}_i) = \min_{w \in \mathcal{C}_i} (\Delta(u_i, w))$

(simplified bound)

Thm: if $0 < \eta_0 < 1 - \sqrt{\rho}$ then $\Pr[\text{Verifier accepts } u_0] \leq \left(1 - \frac{1}{2}\eta_0\right)^m$

if $\eta_0 \geq 1 - \sqrt{\rho}$ then $\Pr[\text{Verifier accepts } u_0] \leq \left(1 - \frac{1}{2}(1 - \sqrt{\rho})\right)^m$

(recall: m is the number of spot checks per round, and $\rho := (\text{rate of } \mathcal{C}_i) = d/|\mathcal{L}|$)

Why is this protocol sound?

Proof idea: To simplify, let's assume that $m = 1$, $\eta_0 < 1 - \sqrt{\rho}$, and

(*) for all $i = 1, \dots, t$ and $r_i \in \mathbb{F}$: $\Delta((u_{i-1})_{\text{fold}, r_i}, \mathcal{C}_i) \geq \Delta(u_{i-1}, \mathcal{C}_{i-1})$

folding does
not decrease
distance

(note: this only holds w.h.p over $r_i \in \mathbb{F}$ by folding corollary)

Then: $\text{Pr}[\text{accept}] = \prod_{i=1}^t \text{Pr}[\text{not reject in round } i] = \prod_{i=1}^t \underbrace{[1 - \Delta(u_i, (u_{i-1})_{\text{fold}, r_i})]}_{\text{prob. } u_i \text{ is accepted after one spot check}}$
↑
independent spot checks per round

$$\leq \exp(-\sum_{i=1}^t \Delta(u_i, (u_{i-1})_{\text{fold}, r_i})) \leq \exp(-\sum_{i=1}^t [\Delta((u_{i-1})_{\text{fold}, r_i}, \mathcal{C}_i) - \Delta(u_i, \mathcal{C}_i)])$$

$\forall x: 1 - x \leq e^{-x} = \exp(-x)$

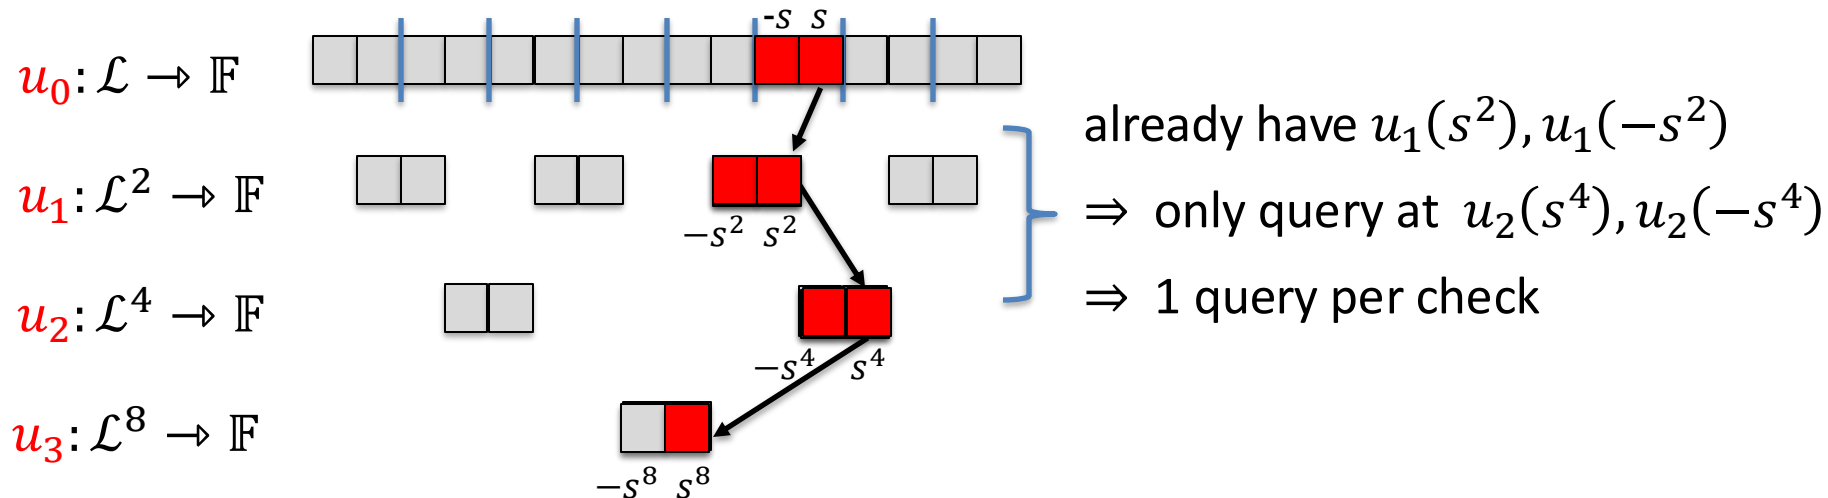
↑
triangular inequality

$\eta_t = 0$ otherwise, Verifier rejects

$$\leq \exp(-\sum_{i=1}^t [\Delta(u_{i-1}, \mathcal{C}_{i-1}) - \Delta(u_i, \mathcal{C}_i)]) = \exp(\eta_t - \eta_0) = \exp(-\eta_0) \leq 1 - \frac{\eta_0}{2}$$

How to spot check: method 2 (the FRI method)

Correlated spot checks: spot check starting at a random $s \in \mathcal{L}$



Total of only mt queries to oracles: m per oracle

(recall: method 1 required $2m$ queries per oracle)

How to spot check: the FRI method


Let $\mathcal{C} = \text{RS}[\mathbb{F}, \mathcal{L}, d]$, $u_0: \mathcal{L} \rightarrow \mathbb{F}$

How to check that $u_i = (u_{i-1})_{\text{fold}, r_i}$:

Repeat m times:

- choose random $s \in \mathcal{L}$, query $(z_0, y_0) \leftarrow (u_0(s), u_0(-s))$
- for $i = 1, \dots, t$:
 - set $s \leftarrow s^2 \in \mathcal{L}^{2^i}$
 - compute $z := (u_{i-1})_{\text{fold}, r_i}(s)$ from z_{i-1}, y_{i-1}
 - query $(z_i, y_i) \leftarrow (u_i(s), u_i(-s))$
 - reject if $z \neq z_i$

only one
query
per round



Why is this sound? (using notation as in earlier proof)

[BKS'18, Thm 7.2]

Proof idea: To simplify, let's assume that $m = 1$, $\eta_0 < 1 - \sqrt{\rho}$, and

(*) for all $i = 1, \dots, t$ and $r_i \in \mathbb{F}$: $\Delta((u_{i-1})_{\text{fold}, r_i}, \mathcal{C}_i) \geq \Delta(u_{i-1}, \mathcal{C}_{i-1})$

folding does
not decrease
distance

Then: $\Pr[\text{reject}] = \Pr[\bigcup_{i=1}^t (\text{not reject in round } i)] = \sum_{i=1}^t \Pr[\text{not reject in round } i]$

can be made into a union of disjoint events

$$= \sum_{i=1}^t \Delta(u_i, (u_{i-1})_{\text{fold}, r_i}) \geq \sum_{i=1}^t [\Delta((u_{i-1})_{\text{fold}, r_i}, \mathcal{C}_i) - \Delta(u_i, \mathcal{C}_i)]$$

↑
triangular inequality

$\eta_t = 0$ otherwise, Verifier rejects

$$\stackrel{(*)}{\geq} \sum_{i=1}^t [\Delta(u_{i-1}, \mathcal{C}_{i-1}) - \Delta(u_i, \mathcal{C}_i)] = \eta_0 - \eta_t = \eta_0 \Rightarrow \Pr[\text{accept}] \leq 1 - \eta_0$$

Comparing methods 1 vs. 2

The FRI method has better soundness: for $\eta_0 = \Delta(u_0, \mathcal{C}_0) < 1 - \sqrt{\rho}$

- **method 1:** $\Pr[\text{Verifier accepts } u_0] \leq \left(1 - \frac{1}{2}\eta_0\right)^m$
- **FRI method:** $\Pr[\text{Verifier accepts } u_0] \leq (1 - \eta_0)^m$ (lower prob. \Rightarrow better bound)

To obtain a SNARK via the BCS'16 compiler we need round-by-round soundness

- **method 1:** independent spot checks \Rightarrow easy to prove R-by-R soundness
- **FRI method:** correlated spot checks \Rightarrow harder to prove R-by-R soundness

(see [[BJKTTZ'23](#)] for R-by-R analysis of FRI)

The number of spot checks m


Goal: u_0 is δ -far from $\text{RS}[\mathbb{F}, \mathcal{L}, d] \Rightarrow \Pr[\text{Verifier accepts } u_0] \leq 1/2^{128}$

For $\delta < 1 - \sqrt{\rho}$: when u_0 is δ -far from $\text{RS}[\mathbb{F}, \mathcal{L}, d]$ we know that

$$\Pr[\text{FRI Verifier accepts } u_0] \leq (1 - \delta)^m$$

So: we want m where $(1 - \delta)^m \leq 1/2^{128}$

$$\Rightarrow m \geq -128 / \log_2(1 - \delta)$$



Main point: (1) the bigger δ is, the smaller m needs to be

(2) smaller $m \Rightarrow$ shorter proof and faster verifier

Choosing the code rate $\rho = d/|\mathcal{L}|$

in practice, set $\delta = \delta_{\max} \approx 1 - \sqrt{\rho}$ (to get smallest possible m)

Example 1: $\rho = 1/2$

$\Rightarrow \delta_{\max} \approx 0.29, \quad |\mathcal{L}| = 2d$

(Plonky3)

Proof length: Longer

Prover work:

Less

shorter codewords \Rightarrow less work for Prover to commit

Example 2: $\rho = 1/4$

$\Rightarrow \delta_{\max} \approx 0.5, \quad |\mathcal{L}| = 4d$

Shorter (smaller m)

More

Is 128-bit security enough??

Suppose m is such that $\Pr[\text{FRI Verifier accepts a } \underline{\text{far}}\ u_0] \leq 2^{-128}$

Fact 1: An adversary that runs FRI 2^{128} times will find a run with favorable spot checks (and forge a proof) with probability $\approx 1/2$

Fact 2: An adversary that runs FRI 2^{80} times will find a run with favorable spot checks (and forge a proof) with probability $\approx 2^{-48}$

For most applications this is sufficient

\Rightarrow do not use less than 120-bits of security; otherwise a 2^{80} adv. will forge proofs.

FRI variants

- (1) Higher-order folding
- (2) Batch FRI for varying degrees
- (3) Reduce proof size by grinding
- (4) STIR and WHIR variants

(1) The benefits of higher-order folding

$$|\mathcal{L}| = n$$


$$u_0: \mathcal{L} \rightarrow \mathbb{F}$$

$$u_1: \mathcal{L}^2 \rightarrow \mathbb{F}$$

$$u_2: \mathcal{L}^4 \rightarrow \mathbb{F}$$

$$u_3: \mathcal{L}^8 \rightarrow \mathbb{F}$$

2-way folding **three times**




Protocol has $t = \log_2 n$ rounds

But: total of 6 queries per 8-way step

Can shrink Merkle proofs by placing entire coset in one leaf of Merkle tree:

\Rightarrow 3 Merkle proofs are about $3\log_2 n$ hashes

8-way folding **once**



Protocol has $\log_8 n = t/3$ rounds

vs. total of 8 queries per 8-way step

\Rightarrow 1 Merkle proof is about $\log_2 n$ hashes

(2) Batch FRI for varying degrees: two methods

Poly-IOPP often do multiple RS proximity tests:

Let $u_j: \mathcal{L}_j \rightarrow \mathbb{F}$ be words for $j = 1, \dots, k$

All u_1, \dots, u_k are encoded with the same rate $\rho = d_j/|\mathcal{L}_j|$

Goal: reject if for any j , u_j is δ -far from $\text{RS}[\mathbb{F}, \mathcal{L}_j, d_j]$



different d_j

The plan: test all k words as a batch.

Method 1: degree padding

Let $d_{\max} := \max_j d_j$, $\mathcal{L}_{\max} := \bigcup_j \mathcal{L}_j$

$u_1: \mathcal{L}_1 \rightarrow \mathbb{F}$ 

$$d_1 = d_{\max} \quad (= \rho \cdot |\mathcal{L}_{\max}|)$$

$u_2: \mathcal{L}_2 \rightarrow \mathbb{F}$ 

$$d_2 = d_{\max}/4$$

$u_3: \mathcal{L}_3 \rightarrow \mathbb{F}$ 

$$d_3 = d_{\max}/2$$


$u_4: \mathcal{L}_4 \rightarrow \mathbb{F}$ 

$$d_4 = d_{\max}$$

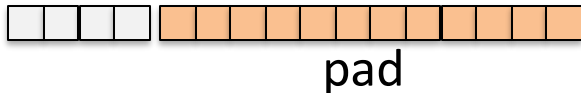
Honest prover can interpolate all u_1, \dots, u_k to \mathcal{L}_{\max}

Method 1: degree padding


Let $d_{\max} := \max_j d_j$, $\mathcal{L}_{\max} := \cup_j \mathcal{L}_j$

$u_1: \mathcal{L}_{\max} \rightarrow \mathbb{F}$ 

$$d_1 = d_{\max} \quad (= \rho \cdot |\mathcal{L}_{\max}|)$$

$u_2: \mathcal{L}_{\max} \rightarrow \mathbb{F}$ 

$$d_2 = d_{\max}/4$$

$u_3: \mathcal{L}_{\max} \rightarrow \mathbb{F}$ 

$$d_3 = d_{\max}/2$$

$u_4: \mathcal{L}_{\max} \rightarrow \mathbb{F}$ 

$$d_4 = d_{\max}$$

Prover can interpolate all u_1, \dots, u_k to \mathcal{L}_{\max}

Method 1: degree padding

Now, batch as follows: Verifier samples a random $r \in \mathbb{F}$, sends to prover

Honest prover defines: $v_r: \mathcal{L}_{\max} \rightarrow \mathbb{F}$ as

$$v_r(a) := \sum_{j=0}^k \sum_{i=0}^{d_{\max}-d_j} r^{e_{i,j}} \cdot [a^i \cdot u_j(a)]$$

where $e_{i,j}$ is a running counter

Example: suppose $(u_1, d_1), (u_2, d_2), (u_3, d_3)$ s.t. $d_1 = 3, d_2 = 4, d_3 = 5$.

$$v_r(a) := \underbrace{[u_1(a) + r \cdot au_1(a) + r^2 \cdot a^2 u_1(a)]}_{\text{linear comb. of } u_1(a), au_1(a), a^2 u_1(a)} + \underbrace{[r^3 \cdot u_2(a) + r^4 \cdot au_2(a)]}_{\text{linear comb. of } u_2(a), au_2(a)} + r^5 \cdot u_3(a)$$

linear comb. of $u_1(a), au_1(a), a^2 u_1(a)$

$$u_1 \in \text{RS}[\mathbb{F}, \mathcal{L}_1, 3] \Rightarrow a^2 u_1(a) \in \text{RS}[\mathbb{F}, \mathcal{L}_1, 5]$$

linear comb. of $u_2(a), au_2(a)$

$$u_2 \in \text{RS}[\mathbb{F}, \mathcal{L}_2, 4] \Rightarrow au_2(a) \in \text{RS}[\mathbb{F}, \mathcal{L}_2, 4]$$

Method 1: degree padding

Lemma: [STIR, Lemma 4.13] the transform $(u_1, \dots, u_k; r) \rightarrow v_r$ is distance preserving

case 1: (the honest case)

if $\forall j: u_j \in \text{RS}[\mathbb{F}, \mathcal{L}_j, d_j]$ then $v_r \in \text{RS}[\mathbb{F}, \mathcal{L}_{\max}, d_{\max}]$ for all r .

case 2: (the dishonest case)

if some u_j is δ -far from $\text{RS}[\mathbb{F}, \mathcal{L}_j, d_j]$ then
 v_r is δ -far from $\text{RS}[\mathbb{F}, \mathcal{L}_{\max}, d_{\max}]$, w.h.p over r .

The proof follows directly from the RS proximity gap (the BCIKS'20 theorem)

Prover now uses an RS-IOPP to prove that v_r is δ -close to $\text{RS}[\mathbb{F}, \mathcal{L}_{\max}, d_{\max}]$

Method 2: pipelining (no padding or interpolation)

Prover $P(\mathcal{C}, (u_1, u_2, u_3), \cdot)$

Verifier $V^{u_1, u_2, u_3}(\mathcal{C})$


Phase 1: (commit)

suppose $d_2 = d_1/2$ and $d_3 = d_1/4$

honest prover:

$$w_1 := (u_1)_{\text{fold}, r_1} + r_1^2 u_2$$

(fold u_2 into w_1)




$$|\mathcal{L}^2| = |\mathcal{L}|/2$$

honest prover:

$$w_2 := (w_1)_{\text{fold}, r_2} + r_2^2 u_3$$

(fold u_3 into w_2)



$$|\mathcal{L}^4| = |\mathcal{L}|/4$$

⋮

⋮

(3) Grinding to reduce # of spot checks

Prover $P((\mathcal{C}, \delta), u_0, \cdot)$

Verifier $V^{u_0}((\mathcal{C}, \delta))$

Commit phase:



Query phase: Derive queries for spot checks by hashing all MerkleCommits (Fiat-Shamir)

Goal: reduce the number of spot checks m (to reduce proof size)

The problem: reducing m below the computed bound enables adversary to try multiple MerkleCommits, until it finds a favorable set of spot checks

Grinding to reduce # of spot checks

One option: add a grinding phase after commit phase (often used in FRI)

Commit phase:



Grind: Find G s.t. $\text{MSB}(\text{SHA3}(G, \text{MerkleCommits}, \text{nonce})) = 0^{64}$

Query phase: Derive queries for spot checks by hashing
(Fiat-Shamir) **all MerkleCommits AND G**

Nonce prevents adversary from pre-computing G (e.g, nonce = head of blockchain)

Why does grinding help?

Adversary: wants δ -far u_0 for which it can generate an RS-proximity proof

FRI without grinding:

m is set so that $E[\text{time to find } \delta\text{-far } u_0 \text{ with favorable queries}] \geq 2^{128}$

\Rightarrow time to find a false proximity proof is $\approx 2^{128}$

FRI with grinding: every u_0 attempt takes time $\approx 2^{64}$ to find G

\Rightarrow suffice that $E[\text{time to find } \delta\text{-far } u_0 \text{ with favorable queries}] \geq 2^{64}$


\Rightarrow can halve the number of spot checks m $\left(128 / -\log_2(1-\delta) \rightarrow 64 / -\log_2(1-\delta) \right)$

\Rightarrow shrink proof length by about $\times 2$

(4) STIR: an FRI variant

[ACFY'24]

Recall: in FRI, distances and # spot checks m are fixed round-to-round




STIR: an FRI variant

[ACFY'24]

STIR main idea: in each round reduce the rate and increase distance


\Rightarrow # spot checks can be decreased from round to round



STIR: an FRI variant

[ACFY'24]

Idea 1: reduce the code rate by making the honest prover interpolate




New goal: prove v_0 is δ -close to $\text{RS}[\mathbb{F}, \omega\mathcal{L}^2, d/4]$

$$\text{rate} = (d/4)/(|\mathcal{L}|/2) = \rho/2 \Rightarrow \text{lower rate}$$

STIR: an FRI variant

[ACFY'24]

Idea 1: reduce the code rate by making the honest prover interpolate



Main point: Lower rate $\rho/2$ means # spot checks for new goal v_0 can be smaller.


Rate drops by a factor of 2 after every folding step \Rightarrow shorter overall proof

STIR: an FRI variant

[ACFY'24]

The problem: now we cannot spot check $v_0: \omega\mathcal{L}^2 \rightarrow \mathbb{F}$

$u_0: \mathcal{L} \rightarrow \mathbb{F}$



send $v_0: \omega\mathcal{L}^2 \rightarrow \mathbb{F}$ to Verifier

Can only spot check here

how to spot check this?

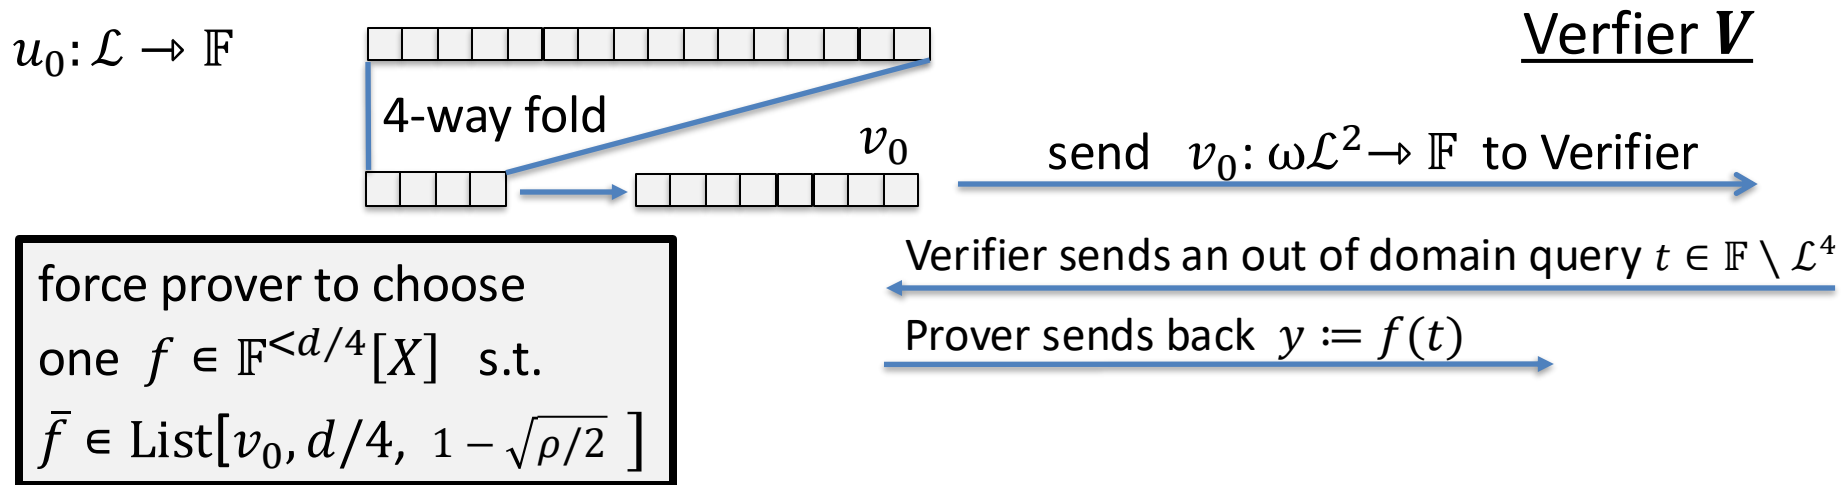
Its 4th roots are not in \mathcal{L} !!

Idea 2: use quotienting for two things:

- (1) spot checks on v_0 , and
- (2) for a malicious prover, increase distance to $\text{RS}[\mathbb{F}, \omega\mathcal{L}^2, d/4]$

How to spot check v_0 by quotienting

How? Honest prover will quotient v_0 by the query points




Honest prover will use $f \in \mathbb{F}^{<d/4}[X]$ s.t. $\bar{f} = (u_0)_{4\text{fold},r}$ on \mathcal{L}^4

How to spot check v_0 by quotienting

How? Honest prover will quotient v_0 by the query points

$u_0: \mathcal{L} \rightarrow \mathbb{F}$



Honest prover defines u_1 as the result of quotienting v_0 by $\{ (t, y), (s_1, y_1), \dots, (s_m, y_m) \}$

$u_1: \omega\mathcal{L}^2 \rightarrow \mathbb{F}$



Verifier V

send $v_0: \omega\mathcal{L}^2 \rightarrow \mathbb{F}$ to Verifier

Verifier sends an out of domain query $t \in \mathbb{F} \setminus \mathcal{L}^4$

Prover sends back $y := f(t)$


Verifier sends spot check points $s_1, \dots, s_m \in \mathcal{L}^4$

Prover sends back $y_i := f(s_i), i = 1, \dots, m$

How to spot check v_0 by quotienting

How? Honest prover will quotient v_0 by the query points

$u_0: \mathcal{L} \rightarrow \mathbb{F}$



That is, $u_1(a) := \frac{v_0(a) - I(a)}{V(a)}$ where

- $I(t) = y$ and $I(s_i) = y_i$
- $V(t) = 0$ and $V(s_i) = 0$

$u_1: \omega\mathcal{L}^2 \rightarrow \mathbb{F}$



Verifier V

send $v_0: \omega\mathcal{L}^2 \rightarrow \mathbb{F}$ to Verifier

Verifier sends an out of domain query $t \in \mathbb{F} \setminus \mathcal{L}^4$

Prover sends back $y := f(t)$

Verifier sends spot check points $s_1, \dots, s_m \in \mathcal{L}^4$


Prover sends back $y_i := f(s_i), i = 1, \dots, m$

In query phase, Verifier computes y_i itself by querying u_0 at $4m$ points and folding

How to spot check v_0 by quotienting

How? Honest prover will quotient v_0 by the query points

$u_0: \mathcal{L} \rightarrow \mathbb{F}$



Verifier V

send $v_0: \omega\mathcal{L}^2 \rightarrow \mathbb{F}$ to Verifier

That is, $u_1(a) := \frac{v_0(a) - I(a)}{V(a)}$ where

- $I(t) = y$ and $I(s_i) = y_i$
- $V(t) = 0$ and $V(s_i) = 0$

$u_1: \omega\mathcal{L}^2 \rightarrow \mathbb{F}$



Verifier sends an out of domain query $t \in \mathbb{F} \setminus \mathcal{L}^4$

Prover sends back $y := f(t)$

Verifier sends spot check points $s_1, \dots, s_m \in \mathcal{L}^4$

Prover sends back $y_i := f(s_i), i = 1, \dots, m$

Iterate to prove that u_1 is $(1 - \sqrt{\rho/2})$ -close to $\text{RS}[\mathbb{F}, \omega\mathcal{L}^2, d/4]$ with a smaller m

STIR: summary

Main benefit: STIR proof is about $2 \times$ shorter than FRI proof
(using the same rate ρ for the input u_0)

Cons:

- Prover is a bit slower because of interpolation and quotienting
- Verifier is a bit slower because of quotienting
- Batching via pipelining is more cumbersome:
 - Often, functions to batch are all defined using the same rate ρ , but STIR iterations use a different rate in every round
 \Rightarrow Prover will need to interpolate functions to expand to lower rate

WHIR: better than STIR [ACFY'24]

WHIR: combines all spot checks into a Sumcheck \Rightarrow no quotienting

- Fold k levels per round, but Verifier now does fewer field ops.
 \Rightarrow fast verifier ($\approx 1.9\text{M}$ gas [in the EVM](#))
- Supports queries to a multilinear polynomial (not just univariate)

How? Not today. (Builds on BaseFold)

The future: other codes

Are there better codes than Reed-Solomon?

The problem with RS-based SNARKs


- **Not field agnostic:** requires an n -th primitive root of unity in \mathbb{F}
 - \Rightarrow can only use fields where $n = |\mathcal{L}|$ divides $|\mathbb{F}| - 1$
 - \Rightarrow difficult to support specific fields (e.g., for ECDSA arithmetic)
- **Encoding is done via an FFT:** takes time $O(n \cdot \log n)$
 - \Rightarrow when $n \approx 2^{20}$, the $(\log_2 n)$ causes $20 \times$ work for prover
- **FRI enables:** a (univariate Poly-IOP) \rightarrow IOP compiler
 - \Rightarrow what about (multilinear Poly-IOP) \rightarrow IOP compiler??
 - e.g., $g(x_1, x_2, x_3) = 5x_1 + 2x_2 + 4x_1x_2 + 12x_1x_3 + 7x_1x_2x_3$

FRI-like proximity proof for other linear codes

FRI can be generalized to any $[n, k, l]_p$ **linear code** $\mathcal{C} \subseteq \mathbb{F}^n$ where:

There is a sequence of linear codes $\mathcal{C} = \mathcal{C}_0, \mathcal{C}_1, \dots, \mathcal{C}_t$ s.t.

1. \mathcal{C}_i is a $[n_i, k_i, l_i]_p$ linear code where $n = n_0 > n_1 > \dots > n_t$ and n_t is “sufficiently small”,
2. there are distance preserving maps $\mathcal{C}_{i-1} \rightarrow \mathcal{C}_i$ for $i = 1, \dots, t$,
3. there is a “fast” encoding algorithm $\mathbb{F}^k \rightarrow \mathcal{C}$, and
4. min-distance of \mathcal{C}_0 is sufficiently large
(to reduce # of spot checks)



A proximity proof for other linear codes

A **field agnostic** proximity test: (e.g., FRI over the ECDSA prime)

- Gives a (univariate Poly-IOPP) \rightarrow IOPP over an arbitrary prime p

(1) ECFFT [[BCKL'22](#)]:

FRI using functions over an elliptic curve E/\mathbb{F}_p ,
where the order of $E(\mathbb{F}_p)$ is divisible by n
(even though p is not)

(2) a proximity proof for algebraic geometric codes [[BLNR'20](#)].
Here polynomials are replaced with “functions on a curve”


A proximity proof for other linear codes

Circle Stark [HLP'24]: Let $M_{11} := 2^{31} - 1$ (the 11th Mersenne prime)

arithmetic mod M_{11} is super fast, ($x \bmod M_{11}$ is just an addition)
but $M_{11} - 1$ is not divisible by a high power of 2

Instead: run FRI over the projective line mod M_{11} whose size is $M_{11} + 1$

One way to represent the projective line is as points on a circle:



divisible by a high
power of 2

So: the circle $(x^2 + y^2 = 1)$ is the same as $\mathbb{F}_p \cup \{\infty\} \Rightarrow (p + 1)$ points

BaseFold [ZCF'23]

BaseFold: generalizes FRI to any foldable code

⇒ The generalization gives a **field agnostic** proximity test

[note: every foldable code is a multilinear Reed-Muller code]

For a (multilinear Poly-IOPP) → IOPP compiler need a multilinear PCS

- The problem: quotienting only applies to univariates
- BaseFold solution:

build a multilinear PCS from Sumcheck and a proximity test

(also adopted into Whir)

How? Not today.

More SNARK-useful linear codes

Spielman codes: [[BCG'20](#), [Breakdown'21](#), [Orion'22](#)]

- Linear codes with a good minimum distance and a very fast (linear time) encoding algorithm $\mathbb{F}^k \rightarrow \mathcal{C}$.
- Also field-agnostic.

Cons: large IOPP proof \Rightarrow large SNARK proof

Expand Accumulate codes: [[BFKTWZ'24](#)]

- Field-agnostic codes, but shorter proofs than Breakdown

Cons: $O(n \cdot \log n)$ time encoding.

More SNARK-useful linear codes

Repeat-Accumulate-Accumulate (RAA) codes: [[Blaze'24](#)]

- Constructs a multilinear polynomial commitment over \mathbb{F}_{2^k} with a fast (linear time) prover time and $O(\log^2 n)$ proof size
 $\Rightarrow \mathbb{F}_{2^k}$ is friendly to modern CPU instructions
- The commitment uses the tensor code approach of [[BCG'20](#)] (making use of Sumcheck).

Much more to do in non-RS based SNARKs

Further reading

- [FRI](#) (2018) and [analysis](#) (2018): Fast Reed–Solomon Interactive Oracle Proofs of Proximity
- [DEEP-FRI](#) (2019): Out of domain sampling improves soundness
- [BCIKS](#) (2020): Proximity Gaps for Reed–Solomon Codes
- [CircleSTARK](#) (2024): FRI using a Mersenne prime
- [STIR](#) (2024): Reed–Solomon proximity testing with fewer queries
- [WHIR](#) (2024): Reed–Solomon proximity testing with a fast verifier

Beyond Reed-Solomon codes (a few recent results):

- [Breakdown](#) (2021), [Orion](#) (2022): Polynomial commitments with a fast prover
- [BaseFold](#) (2023): Efficient Polynomial commitments from foldable codes
- [Blaze](#) (2024): Fast SNARKs from Interleaved RAA Codes

END OF MODULE